Supplementary material

A Model parameterization and algebraic properties

The frozen ESM-2 model supplies 1,280-dimensional residue representations. Independent linear projections map wild-type, local-change, and cross-site-change features to the selected projection dimension. Amino-acid states use one-hot encodings. The single and pair heads have a 256-dimensional hidden layer, a 128-dimensional second layer, GELU activations, and a scalar output. The pair head also receives the calibrated single effects and their observation indicators. Its sequence relation uses the logarithm of one plus the site gap and the gap divided by sequence length. Projection dimensions, learning rates, auxiliary weights, and ω are selected from source-validation measurements.

The cross-site context is the change at one site after the other site is substituted,

$$\chi_{i\leftarrow j}(b) = h_i(x_p(j,b)) - h_i(s_p). \tag{11}$$

Each single-mutant sequence is encoded once. Its residue states are reused across candidate pairs. The pair network therefore receives mutation-dependent context without encoding every double-mutant sequence. Sequence separation supplies the relation between endpoints. The structural variant adds reference Cα distances and residue confidence.

The normalized phenotype map is

$$g(z) = \frac{z + \sum_{k=1}^{20} b_k[\tanh(c_k z + d_k) - \tanh(d_k)]}{1 + \sum_{k=1}^{20} b_k c_k \operatorname{sech}^2(d_k)}, \qquad b_k, c_k > 0. \tag{12}$$

Positive parameters use softplus. Substitution of $z = 0$ gives $g(0) = 0$. Differentiation gives $g'(0) = 1$. Its derivative is positive for every z, and its linear term preserves nonzero slopes in the tails. The positive factor τ sets the phenotype scale.

For $u = \widetilde{a}_i(a)$ and $v = \widetilde{a}_j(b)$, the predicted phenotype contrast decomposes as

$$\widehat{\epsilon}_{\text{global}} = \tau[g(u+v) - g(u) - g(v)], \tag{13}$$

$$\widehat{\epsilon}_{\text{pair}} = \tau[g(u+v+r_{ij}) - g(u+v)]. \tag{14}$$

Their sum is $\widehat{\epsilon}_p$. With the identity map and unit scale, the phenotype contrast equals r_{ij}. With a nonlinear map, the complete contrast includes both contributions.

Symmetry of F implies symmetry of its four-state contrast. Setting $b = w_j$ makes its first two terms equal and its last two terms equal, giving $r_{ij}(a, w_j) = 0$. The other reference constraint follows identically. For an endpoint term $v_i(a)$, the contrast contains $v_i(a) - v_i(a) - v_i(w_i) + v_i(w_i) = 0$. The second endpoint term and the constant cancel in the same way. These are algebraic properties of the parameterization. Real-data model checks also verify endpoint exchange and reference-state outputs.

Matched controls for the pair function

The input-removal control supplies model-predicted single effects $a_i(a), a_j(b)$ to every evaluation of F and sets its observation indicators to zero. Equation (3) continues to define the external single-effect sum. Measured values and their availability indicators therefore reach that sum through the original permission rule. The pair network receives neither channel. Its input dimensions and parameter count remain unchanged.

The direct-head control predicts $r_{ij}(a,b) = F(i,a;j,b)$ with the same symmetric network. Its projections, endpoint features, measured inputs, hidden dimensions, and parameter count match the four-state model. This intervention isolates the reference contrast while retaining endpoint symmetry and the external single-effect calibration.

Both controls reuse the full model's source-selected hyperparameters in each outer partition. They use the same pretrained single checkpoint for each training repetition, random initialization, source rows, sampler, optimizer, objectives, and checkpoint-selection rule. The maximum epoch budget and early-stopping patience are identical. Realized stopping epochs and fitted weights may differ. The label-mixing weight follows the same source-validation procedure. Full-target evaluation retains the same proteins, candidate pools, measurement margins, budgets, and aggregation. The stricter homology grouping repeats this comparison independently. These controls separate the predictive effects of the four-state construction and measured pair inputs from their algebraic definitions.

Table 4. Paired native top-10 hit-rate advantage of SCEPI. Gains and 95% bootstrap intervals are in percentage points. Positive fractions refer to resampled mean gains. Every comparison includes 120 proteins, 44 groups, and five model repetitions.

Comparator	Mean gain	95% interval	Positive fraction
Additive	4.83	[1.97, 7.95]	0.9991
Global nonlinear	6.77	[3.10, 10.11]	1.0000
ThermoMPNN-D	4.15	[−0.32, 7.08]	0.9628
Kermut	5.82	[2.59, 8.79]	0.9996
StabilityArc + Kermut	2.47	[−1.17, 5.55]	0.9165
Mean MLP	1.15	[−2.90, 3.76]	0.7387
Mean ridge	−0.20	[−4.10, 1.92]	0.4312

Table 5. Paired top-10 hit-rate advantage under common measured main effects at quartet margin $\geq$ 0.5 kcal/mol. Gains and 95% bootstrap intervals are in percentage points. Every comparison includes 115 proteins, 43 groups, and five model repetitions.

Comparator	Mean gain	95% interval
Additive	6.63	[2.79, 9.96]
Global nonlinear	6.33	[2.38, 9.39]
ThermoMPNN-D	5.88	[2.39, 8.55]
Kermut	5.86	[2.19, 9.15]
StabilityArc + Kermut	3.03	[−0.64, 6.41]
Mean MLP	3.84	[0.27, 5.76]
Mean ridge	3.67	[0.77, 6.13]

B Repetitions, aggregation, and statistical analysis

The 30% grouping has five outer partitions. Each source-trained model uses one partition for testing, the next partition for source validation, and the remaining three for source training. The main comparisons contain all 25 outer-fold and training-seed evaluations per model. Training seeds are 11, 23, 47, 89, and 131. Every eligible target protein is evaluated once per training seed in its assigned outer test partition. All eligible target double mutants enter the full-target evaluation. There are 795 paired protein-repetition records for phenotype prediction, 775 for residual prediction, and 600 for standard compensation selection. The 0.5 kcal/mol margin cohort contains 575 records. Each protein is averaged over its five repetitions before the equal-protein mean is computed.

At each repetition, all compared methods use the same permitted target labels and candidate definition. Means use the paired cohort. Phenotype and residual comparisons cover 159 and 155 proteins, respectively, from 51 homology groups. Standard compensation covers 120 proteins and 44 groups. The 0.5 kcal/mol margin comparison covers 115 proteins and 43 groups. Source-validation configurations are retained across training repetitions. Candidate yield is excluded from source hyperparameter selection. Five outer partitions describe source-data separation; five training seeds describe repeated model fitting.

Uncertainty uses 10,000 homology-group bootstrap draws with statistical seed 20261013. Each draw samples groups with replacement and retains all member proteins. Its effect is the equal-protein mean over the sampled groups. Repeated model fits are averaged within protein before resampling. Comparisons with the same group membership use the same resampling draws. A bootstrap positive fraction is the fraction of resampled mean differences above zero. It describes uncertainty in the observed paired mean under this resampling scheme. Tables 4 and 5 report the top-10 endpoint. Table 6 reports the matched pair-function controls.

Measured pair inputs improve both mean residual error and native hit rate in all five repetitions under each grouping. Their residual comparison has 143 protein-level wins under the primary grouping and 142 under the stricter grouping, out of 155 proteins. The reference contrast lowers mean residual error in all five

Table 6. Paired advantage of SCEPI over matched pair-function controls under the 30% and 25% identity groupings. Residual RMSE reductions are in kcal/mol; native top-10 hit-rate gains are in percentage points. Brackets give 95% homology-group bootstrap intervals. Residual and selection comparisons contain 155 and 120 proteins, respectively. Each protein has five matched repetitions.

Grouping	Pair-function control	Residual reduction	Hit-rate gain
30%	Predicted singles in F	0.172 [0.152, 0.211]	3.03 [0.72, 4.32]
30%	Direct residual	0.041 [−0.014, 0.073]	0.52 [−3.04, 2.46]
25%	Predicted singles in F	0.210 [0.155, 0.263]	3.00 [−0.06, 4.36]
25%	Direct residual	0.026 [−0.019, 0.042]	−0.23 [−1.58, 0.81]

repetitions at each threshold, with 106 and 97 protein-level wins. Its native hit-rate gain is positive in four repetitions under the 30% grouping and two under the 25% grouping.

C Source model selection and implementation

Complete source quartets provide the residual target in Equation (1). For source proteins P, the residual objective gives equal weight to proteins and site pairs,

$$\mathcal{L}_\epsilon = \frac{1}{|P|}\sum_{p\in P}\frac{1}{|J_p|}\sum_{(i,j)\in J_p}\frac{1}{|Q_{pij}|}\sum_{q\in Q_{pij}}\mathrm{Huber}\left(\frac{\widehat{\epsilon}_q-\epsilon_q}{s}\right), \tag{15}$$

where J_p contains measured site pairs and Q_{pij} contains their quartets. The stability scale is $s = 1$ kcal/mol. Sampling across proteins, then pairs, then quartets implements this weighting.

The fitted objective is $\mathcal{L}_\epsilon + \lambda_y \mathcal{L}_y$. The observation term $\mathcal{L}_y$ is the negative log-likelihood for the predicted phenotype. Point observations use a Gaussian density; interval observations use the corresponding Gaussian interval probability. Each original observation appears once in this term. This criterion jointly fits the single head, pair function, and phenotype map. Residual-only and phenotype-only models are trained separately.

The single-effect model is fitted to source single-mutant measurements before joint training. The joint configuration grid contains projection dimensions of 64 and 128, learning rates of 10^{-4}, 3×10^{-4}, and 10^{-3}, and auxiliary phenotype weights of 0.1 and 1. The resulting twelve configurations are fitted separately in each outer partition. AdamW uses weight decay 10^{-4} and batches of 1,024 records. Training runs for at most 100 epochs, with early stopping after ten epochs without validation improvement. The residual loss uses Huber threshold 1 on the kcal/mol scale.

Training samples three single-label availability conditions. With probability 0.5, it hides target single labels. With probability 0.25, it permits 128 single labels. With probability 0.25, it permits all single labels. Source validation selects $\omega \in \{0, 0.5, 1\}$ in the 128-label condition. It selects the model configuration by a weighted residual RMSE: 0.5 for zero labels, 0.25 for 128 labels, and 0.25 for all single labels. The selected configuration is then refitted for each training repetition. Phenotype-only, residual-only, and representation variants have separate fitted parameters. The additive model uses phenotype validation because its four-term residual is zero by construction.

Frozen sequence representations are cached by reference and single-mutant sequence. The pair predictor reuses those caches across double mutants. The structural variant adds Cα distance and residue confidence from the reference structure. Geometry permutations preserve the recorded distance values while changing their pair assignments. Inference interventions use the fitted structural model. Retrained interventions also use the changed assignments during fitting. Each comparison retains the same outer partition and information condition.

ThermoMPNN-D is retrained on source proteins with its double-mutant head and over-and-back augmentation. Its additive counterpart is trained separately. Kermut fits the permitted target labels with the full composite kernel. Mean ridge and mean MLP fit those target labels using mean-pooled residue representations. ESM-2 additive and context scores provide fixed pretrained representation controls. The main external comparison evaluates each method's native predictions; the common-main-effect comparison replaces its single-effect sum with the same measured sum.

D StabilityArc architecture and target calibration

The StabilityArc comparison uses frozen ESMC-600M wild-type residue embeddings and the author's symmetrized, APC-corrected, entropy-weighted attention contacts [13]. Its decoder has four layers, hidden dimension 256, and eight attention heads with rotary position encoding. The full amino-acid embedding, sequence-distance features, and symmetric contact-aware pair residual are retained. Each of the ten source partitions is trained with the five fixed training seeds, giving 50 source models. Every source model excludes its target homology groups from both fitting and validation.

Each source model is trained for 40 epochs with AdamW, learning rate 10^{-4}, and weight decay 10^{-3}. Each assay contributes at most 40,000 measured variants per epoch. The objective combines Huber loss with 1,024 measured ranking pairs per assay. The Huber threshold is 1, and the two losses have equal weights. Double-mutant regression targets use the source-assay single-mutant median and median absolute deviation. Ranking uses measured stability order. Checkpoints are selected by source-validation double-mutant Spearman correlation. The public release supplies the architecture and a pooled checkpoint; the matched source-partition models are retrained here. Equal loss weights are an explicit implementation choice because the release does not specify them.

Target calibration uses the full Kermut covariance and Gaussian likelihood [6]. The retrained StabilityArc score replaces the prior mean input. All remaining kernel inputs and permitted target labels match the Kermut comparison. Target fitting uses 150 AdamW steps with learning rate 0.1 and weight decay 0.01. Only the permitted reference and single-mutant point labels are supplied. The posterior scores every eligible target double mutant and its constituent singles. The model's four-prediction contrast defines its residual. All 50 source fits and target-prediction groups complete before this comparator enters the reported tables.

E Component comparison details

Table 7. Top-10 compensation hit rates under a common measured single-mutant main-effect sum. Each method contributes its own predicted four-term residual. The standard and margin cohorts contain 120 and 115 proteins, respectively.

Residual source	Standard pool (%)	Margin ≥ 0.5 (%)
SCEPI	**11.27**	**11.15**
ThermoMPNN-D	5.30	5.27
Kermut	5.47	5.29
StabilityArc + Kermut	8.17	8.12
Global nonlinear	4.70	4.82
Additive	4.42	4.52
Mean MLP	7.17	7.30
Mean ridge	7.33	7.48

Inference-time pair removal keeps trained main effects and the map fixed. Retrained additive and global nonlinear models optimize their own parameters. Representation variants retain the same reference-constrained pair construction. The identity-only variant adjusts its hidden width to approximate the main model's trainable parameter count. These comparisons measure component and objective choices within the shared-model class.

Removing cross-site context yields 0.575 kcal/mol residual RMSE and an 11.47% native hit rate under the primary grouping. The amino-acid-identity version reaches 0.557 kcal/mol and 9.12%. Phenotype-only training reaches 0.562 kcal/mol and 11.95%. Under the stricter grouping, phenotype-only training lowers residual RMSE from 0.604 to 0.584 kcal/mol and raises hit rate from 11.15% to 11.98%. The identity phenotype map gives 0.593 kcal/mol and 11.13% under the primary grouping. Cross-site context, the flexible map, and the residual objective remain model and fitting choices with these individual results.

Table 8. Full-target S-all results for completed predictors and model variants. RMSE and recovery are in kcal/mol. Hit rates are native top-10 compensation percentages. The phenotype, residual, and selection cohorts contain 159, 155, and 120 proteins. Each row completes five outer partitions and five training repetitions.

Predictor or variant	Phenotype	Residual	Hit rate	Recovery
SCEPI	0.657	0.595	11.22	−0.257
Additive	0.893	1.091	6.38	−0.387
Global nonlinear	0.911	0.635	4.45	−0.437
ThermoMPNN-D	0.777	0.805	7.07	−0.440
ThermoMPNN additive	0.967	1.091	5.30	−0.511
Kermut	0.798	0.847	5.40	−0.458
StabilityArc + Kermut	0.815	0.853	8.75	−0.344
Mean MLP	1.405	0.966	10.07	−0.413
Mean ridge	1.328	1.002	11.42	−0.354
ESM-2 additive	1.173	1.091	2.50	−0.794
ESM-2 context	1.107	0.864	8.50	−0.499
Pair output removed at inference	1.151	1.072	4.42	−0.432
Direct residual head	0.700	0.635	10.70	−0.302
Measured singles removed from F	0.810	0.767	8.18	−0.396
Cross-site context removed	0.637	0.575	11.47	−0.258
Wild-type context only	0.643	0.573	10.53	−0.273
Amino-acid identities and singles	0.617	0.557	9.12	−0.290
Phenotype-only training	0.619	0.562	11.95	−0.229
Residual-only training	1.071	0.686	8.72	−0.334
Frozen single-effect head	0.689	0.619	9.75	−0.282
Identity phenotype map	0.651	0.593	11.13	−0.257
Structural variant	0.644	0.583	10.82	−0.263
Distance permutation, inference	0.648	0.586	10.95	−0.261
Distance permutation, retrained	0.650	0.590	10.53	−0.270
Distance rewiring, inference	0.646	0.585	10.78	−0.264
Distance rewiring, retrained	0.633	0.572	10.47	−0.257

F Sensitivity to the homology threshold

Grouping at 25% identity with the same coverage requirement produces 45 groups in the phenotype and residual cohorts. The compensation cohort contains 39 groups, and its 0.5 kcal/mol margin subset contains 38. Protein membership remains unchanged. All completed predictors and components repeat the full five-partition, five-training-seed evaluation under this grouping (Table 9). SCEPI's residual RMSE is 0.604 kcal/mol, compared with 0.771 for ThermoMPNN-D, 0.847 for Kermut, 0.874 for StabilityArc + Kermut, and 0.650 for global nonlinearity. Its native compensation hit rate is 11.15%, compared with 6.97%, 5.40%, and 8.23% for the three external methods. Its mean selected recovery exceeds StabilityArc + Kermut by 0.092 kcal/mol.

Under common measured main effects, top-10 hit rates are 11.17% for SCEPI, 4.42% for additivity, 4.63% for global nonlinearity, 5.68% for ThermoMPNN-D, 5.47% for Kermut, 8.57% for mean MLP, and 5.83% for mean ridge. At margin 0.5 kcal/mol, the corresponding rates are 11.23%, 4.52%, 4.75%, 5.65%, 5.29%, 8.77%, and 5.83%. The stricter grouping preserves the average compensation advantage and the common-main-effect comparison.

StabilityArc + Kermut reaches an 8.02% common-main-effect hit rate and 7.95% at margin 0.5 kcal/mol. The paired comparison retains the same proteins and label permissions. Thus the additional source-trained comparator preserves the interaction-selection advantage under both homology thresholds.

The matched direct residual head reaches 0.630 kcal/mol residual RMSE and an 11.38% native hit rate. SCEPI lowers residual error to 0.604 kcal/mol and has an 11.15% native hit rate. Removing measured pair

inputs while retaining the external calibration gives 0.814 kcal/mol and 8.15%. Common-main-effect hit rates are 11.45% for the direct head and 8.15% for the removed-input control. SCEPI reaches 11.17%. These controls retain the residual-accuracy contribution of the reference contrast and the prediction and selection contribution of measured pair inputs.

Table 9. Full-target S-all sensitivity comparison at 25% sequence identity. The columns use the same units and protein cohorts as Table 8. Every row completes five outer partitions and five training repetitions.

Predictor or variant	Phenotype	Residual	Hit rate	Recovery
SCEPI	0.677	0.604	11.15	−0.256
Additive	0.994	1.091	6.65	−0.393
Global nonlinear	0.938	0.650	4.42	−0.432
ThermoMPNN-D	0.758	0.771	6.97	−0.432
ThermoMPNN additive	0.986	1.091	4.95	−0.536
Kermut	0.798	0.847	5.40	−0.458
StabilityArc + Kermut	0.840	0.874	8.23	−0.348
Mean MLP	1.211	0.922	11.20	−0.400
Mean ridge	1.638	1.039	11.25	−0.393
ESM-2 additive	1.173	1.091	2.50	−0.794
ESM-2 context	1.104	0.864	8.50	−0.499
Pair output removed at inference	1.151	1.073	4.42	−0.432
Direct residual head	0.699	0.630	11.38	−0.257
Measured singles removed from F	0.855	0.814	8.15	−0.386
Cross-site context removed	0.664	0.587	11.07	−0.266
Wild-type context only	0.655	0.588	10.12	−0.294
Amino-acid identities and singles	0.622	0.554	9.82	−0.287
Phenotype-only training	0.653	0.584	11.98	−0.235
Residual-only training	1.166	0.682	9.77	−0.311
Frozen single-effect head	0.684	0.611	11.32	−0.256
Identity phenotype map	0.680	0.609	11.07	−0.252
Structural variant	0.683	0.609	11.75	−0.240
Distance permutation, inference	0.685	0.611	11.93	−0.241
Distance permutation, retrained	0.677	0.606	11.75	−0.238
Distance rewiring, inference	0.684	0.609	11.62	−0.243
Distance rewiring, retrained	0.688	0.615	10.62	−0.257

G Evaluation on the designated site-pair subset

The separate site-pair evaluation uses the original unseen-pair test subset within each target protein. It intersects this subset with the same held-out homology groups. Site-pair split seeds are 11, 23, 47, 89, and 131. Candidate eligibility can change across these splits. The completed residual comparison contains 70 proteins, 34 groups, and 339 paired protein-repetition records. Compensation contains 63 proteins, 32 groups, and 269 records. Table 10 reports this subset without combining its values with the full-target results.

Common-main-effect top-10 hit rates on this subset are 5.78% for SCEPI, 3.29% for ThermoMPNN-D, 3.07% for Kermut, 3.91% for global nonlinearity, and 3.62% for additivity. At margin 0.5 kcal/mol, the corresponding rates are 5.38%, 3.12%, 2.67%, 3.64%, and 3.34%. That margin cohort contains 62 proteins, 32 groups, and 263 paired protein-repetition records. Homology-group resampling with 10,000 draws and seed 20261010 gives an additive-comparison gain of 2.04 percentage points and a 95% interval of [0.67, 3.31]. Its gain over global nonlinearity is 1.74 points with an interval of [0.49, 3.12]. In S-all, every target double label is withheld from fitting. The site-pair subset therefore defines an additional evaluation population. The

Table 10. Repeated S-all evaluation on the designated unseen-site-pair subset within held-out homology groups. Each eligible protein is averaged over its available repetitions.

Predictor	Residual RMSE	Native hit rate (%)	Recovery
SCEPI	0.549	5.75	−0.442
ThermoMPNN-D	0.778	4.88	−0.511
Kermut	0.833	3.29	−0.562
Global nonlinear	0.590	3.56	−0.528
Additive	1.080	4.71	−0.493

target-combination-adaptation experiments use the pair split to test extrapolation from permitted training combinations.

H Target-label budgets and unseen-site-pair prediction

The transferred predictor benefits from partial target single-mutant measurements. Across the same 155 residual-evaluation proteins, its RMSE is 0.793 kcal/mol in Z, 0.788 in S32, 0.772 in S128, and 0.595 in S-all. Table 11 reports these four information conditions. Tables 13–15 report the separate target combination-adaptation experiments.

The family-level comparison varies only the target information condition. Source training and validation remain outside each target homology group. All eligible target double mutants enter evaluation. Each condition retains 155 residual-evaluation proteins and 775 protein-repetition records. Kermut requires target measurements, so the Z comparison lists source-trained predictors. Table 11 retains the original four information conditions.

Table 11. Residual RMSE in kcal/mol at different target single-mutant budgets under 30% homology grouping. Each row covers the same 155 proteins and five training repetitions. Z supplies no target labels; S-all supplies all permitted target single-mutant labels.

Condition	SCEPI	Additive	Global nonlinear	ThermoMPNN-D	Kermut
Z	0.793	1.091	0.861	1.006	—
S32	0.788	1.091	0.857	0.943	0.998
S128	0.772	1.091	0.841	0.908	0.928
S-all	0.595	1.091	0.635	0.805	0.847

The available ESM-MSR comparison uses the epoch-10 source-validation checkpoint from family30, fold 0, and repetition seed 11. We compare every method on its matched held-out proteins in that fold. Residual evaluation retains 61 proteins; phenotype evaluation retains 65. Table 12 reports Z and S-all separately. This checkpoint supplies one fold and one repetition of the original matrix. Its results remain separate from the complete five-fold, five-repetition comparison.

The C comparison fits the SCEPI adaptation variant to permitted target training combinations. Its source initialization uses an identity phenotype map and a single-effect head held fixed during source pair fitting. Target adaptation retains that head and updates the pair network and observation noise. Target validation selects the stopping point. The target-only variant fits its single-effect head without supervised source weights. Both variants retain the same target single-label permissions. Pair removal sets the fitted pair output to zero at inference. The tables label the two versions as SCEPI adaptation, source and SCEPI adaptation, target only. These configurations are evaluated separately from the jointly trained S-all predictor.

Each budget includes proteins with enough permitted training combinations and eligible test quartets. Every comparator in a table uses the same protein-repetition intersection. Unseen-site-pair evaluation excludes all test pairs from target combination fitting. Random splitting shares site pairs between training

Table 12. Matched comparison with the available epoch-10 ESM-MSR checkpoint. RMSE uses kcal/mol. Columns report residual RMSE, phenotype RMSE, and phenotype Spearman correlation. Every method shares the same 61 residual proteins and 65 phenotype proteins within each information condition. Kermut requires target labels and is omitted in Z.

	Z			S-all		
Predictor	Residual	Phenotype	Spearman	Residual	Phenotype	Spearman
SCEPI	**0.744**	**1.071**	0.598	**0.550**	**0.621**	**0.767**
ESM-MSR	0.899	1.185	**0.646**	0.891	1.091	0.678
ThermoMPNN-D	1.031	1.237	0.587	0.805	0.772	0.707
Global nonlinear	0.829	1.100	0.534	0.632	0.888	0.735
Additive	1.035	1.136	0.588	1.035	0.885	0.731
Kermut	—	—	—	0.829	0.793	0.731

and test combinations. The auxiliary two-pair cohort retains its designated pair split. Coverage reports proteins and matched protein-repetition records. Metrics average repetitions within protein before equal-protein averaging.

Table 13. Unseen-site-pair residual RMSE in kcal/mol after target combination adaptation. Coverage lists proteins and matched protein-repetition records. All listed methods use the same cohort at each budget. Source initialization and target-only fitting are distinguished. C-all uses all permitted training combinations. Test site pairs remain excluded from fitting.

Predictor	C128	C512	C2048	C-all
Coverage	70 / 339	45 / 219	4 / 17	70 / 339
SCEPI adaptation, source	0.567	0.570	0.424	0.559
Pair removed, source initialization	1.080	1.163	0.967	1.080
SCEPI adaptation, target only	0.799	0.788	0.495	0.774
Pair removed, target only	1.080	1.163	0.967	1.080
SCEPI structure, source	**0.558**	**0.563**	**0.408**	**0.551**
SCEPI structure, target only	0.795	0.785	0.467	0.772
ThermoMPNN-D	0.775	0.787	0.618	0.771
ThermoMPNN additive	1.080	1.163	0.967	1.080
Kermut	0.758	0.770	0.577	0.740
MoCHI, order 1	0.857	0.946	0.834	0.900
MoCHI, order 2	0.859	0.935	0.845	0.895
Mean ridge, target fit	0.933	0.994	0.767	0.921
Mean MLP, target fit	0.666	0.635	0.430	0.620
Mean ridge, source prior	1.000	1.067	0.917	0.999
Mean MLP, source prior	0.849	0.882	0.670	0.816

I Functional assays and higher-order target adaptation

The functional double-mutant cohort contains 74,598 double mutants from six ProteinGym assays [15]. Measurements retain their original scales. The three KRAS abundance blocks remain separate measurement conditions. The resulting eight conditions contain 7,898 single mutants. Seven conditions supply measured wild-type values and 70,198 complete quartets. HIS7 contributes phenotype-ranking tests and has no observed

Table 14. Random-split residual RMSE in kcal/mol after target combination adaptation. The cohort and averaging rule are shared by every row within each budget. Random splitting evaluates new amino-acid combinations at site pairs that may appear in training.

Predictor	C128	C512	C2048	C-all
Coverage	125 / 601	90 / 450	11 / 55	125 / 601
SCEPI adaptation, source	0.355	0.315	0.279	0.277
Pair removed, source initialization	1.065	1.122	1.212	1.065
SCEPI adaptation, target only	0.401	0.358	0.311	0.311
Pair removed, target only	1.065	1.122	1.212	1.065
SCEPI structure, source	**0.351**	**0.310**	**0.273**	**0.274**
SCEPI structure, target only	0.399	0.350	0.303	0.304
ThermoMPNN-D	0.424	0.394	0.381	0.385
ThermoMPNN additive	1.065	1.122	1.212	1.065
Kermut	0.447	0.405	0.355	0.346
MoCHI, order 1	0.737	0.802	0.991	0.810
MoCHI, order 2	0.736	0.800	0.997	0.813
Mean ridge, target fit	0.527	0.498	0.568	0.449
Mean MLP, target fit	0.389	0.339	0.283	0.305
Mean ridge, source prior	0.971	1.031	1.124	0.967
Mean MLP, source prior	0.749	0.746	0.726	0.690

Table 15. Residual RMSE in kcal/mol on the auxiliary two-pair cohort. Coverage lists proteins and matched protein-repetition records. This protocol contains the source-prior and pair-network comparisons shown here.

Predictor	C128	C-all
Coverage	27 / 131	29 / 140
SCEPI adaptation, source	0.683	0.687
Pair removed, source initialization	1.163	1.172
SCEPI adaptation, target only	1.138	1.149
Pair removed, target only	1.163	1.172
SCEPI structure, source	**0.653**	**0.659**
SCEPI structure, target only	1.147	1.157
ThermoMPNN-D	0.813	0.823
ThermoMPNN additive	1.163	1.172
Kermut	0.840	0.827
Mean ridge, source prior	1.063	1.074
Mean MLP, source prior	1.002	1.001

wild type for quartet evaluation. Table 16 lists the full coverage. Each fit uses only its permitted measurement condition.

We evaluate random combination splits and held-out site pairs at C128, C512, C2048, and C-all. Each condition supplies its permitted single mutants and training combinations. Test combinations remain excluded from fitting and validation. The C2048 condition excludes HIS7 because its training pool is smaller than the budget. Phenotype rank correlation is computed on measured test double mutants. Residual evaluation requires at least 30 complete test quartets. To compare assay scales, residual RMSE is divided by the population standard deviation of permitted training point values. Tables report equal-assay means. Repetitions are averaged within condition, the three KRAS conditions are averaged within their assay, and assays receive equal weight.

The functional SCEPI adaptation variant first fits single-effect predictions and a target-specific monotone phenotype map. It then freezes both and fits the pair function. The source-initialized version starts from

Table 16. Accepted functional double-mutant data and complete measured quartets. Each row retains its original measurement condition. The three KRAS conditions belong to one abundance assay. HIS7 lacks an observed wild type and contributes no complete quartets.

Assay condition	Singles	Doubles	Complete quartets
D7PM05_CLYGR fluorescence	1,169	10,148	9,517
GFP_AEQVI fluorescence	1,084	12,777	11,954
HIS7_YEAST selection fitness	168	1,475	0
Q6WV12_9MAXI fluorescence	1,141	15,992	15,717
Q8WTC7_9CNID fluorescence	1,201	11,260	10,822
KRAS abundance, block 1	1,188	11,946	11,904
KRAS abundance, block 2	1,122	7,345	7,184
KRAS abundance, block 3	825	3,655	3,100
Total	7,898	74,598	70,198

source models trained with an identity phenotype map and a fixed single-effect head. Related source proteins are excluded using the recorded homology relation. Eligible source models are averaged at the effect level. The target-only version uses no supervised source weights. These variants use predicted single effects within the pair network. They evaluate source initialization and target combination fitting separately from measured-input conditioning in the primary S-all predictor.

The independent no-pair control fits the first stage without adding pair effects. Inference-time pair removal uses the fitted two-stage model and sets its pair output to zero. Learning rates are 10^{-4}, 3×10^{-4}, and 10^{-3}. Target-only models compare projection dimensions 64 and 128. Pair fitting compares observation-loss weights 0.1 and 1. Each stage runs for at most 100 epochs with ten-epoch validation patience. Target validation selects phenotype rank correlation, with identifiable phenotype RMSE breaking ties.

MoCHI and MAVE-NN compare first-order and second-order sequence models [4, 10]. RFA uses second-order effects with the author's sigmoid phenotype link [14]. Its fixed configuration sets $\lambda = 0$. New optimizations use at most 500 L-BFGS iterations, gradient threshold 10^{-5}, and relative objective tolerance 10^{-4} over 25 iterations. Normalization has a separate 500-iteration cap. Existing checkpoints with the same default penalty are reused without further optimization. Retained validation-selected predictions use the original twelve penalties, $\lambda/n \in \{0, 10^{-6}, 10^{-5.5}, \ldots, 10^{-1}\}$. Their fitted configurations and stopping records remain associated with the reported predictions. No test outcomes enter fitting or configuration selection.

The RFA records contain 350 eligible fits, comprising 310 functional-assay fits and 40 FLIP2 fits. The fixed default supplies 284 results through 264 new fits and 20 reused default checkpoints. The remaining 66 results retain their earlier validation-selected configurations. Thirty additional task records lack the required budget or target training and validation labels. Their exclusion is recorded separately from fitted results.

RFA gives constant GFP predictions in all five C128 repetitions under each split and in two unseen-pair C512 repetitions. Their twelve Spearman correlations are undefined. Rank comparisons use the same finite protein-repetition records for every method. C128 therefore excludes GFP and retains five assays, seven conditions, and 35 records. Unseen-pair C512 retains six assays, eight conditions, and 38 records. The other six-assay rank columns contain 40 records. C2048 retains five assays and 35 records after excluding HIS7. Residual evaluation retains all seven eligible conditions and 35 records at every budget, including the constant-prediction cases.

Tables 17 and 18 separate residual prediction from phenotype ranking. At C128 on unseen pairs, the source-initialized interaction variant has normalized residual RMSE 0.508. Its pair-removed control has 0.551, and the target-only interaction variant has 0.537. First-order MAVE-NN reaches 0.484. At C-all, RFA reaches 0.409 and the source-initialized variant reaches 0.448. On the full six-assay unseen-pair ranking cohort at C128, the source-initialized variant has Spearman correlation 0.730, compared with 0.741 after pair removal. The pair fit improves residual accuracy in this setting while the phenotype-ranking comparison favors the no-pair control.

FLIP2 evaluates TrpB one-to-many, two-to-many, and by-position splits; amylase one-to-many; IRED two-to-many; Hydro three-to-many; and three Hydro background-transfer splits [16]. We retain the continu-

Table 17. Normalized residual RMSE on functional double mutants. Every column contains five assays, seven measurement conditions, and five repetitions. Errors are divided by each condition's training phenotype standard deviation before equal-assay averaging. C-all uses the permitted training combinations. Lower values are better.

	Random combinations				Unseen site pairs			
Predictor	C128	C512	C2048	C-all	C128	C512	C2048	C-all
SCEPI adaptation, source	0.503	0.484	0.465	0.435	0.508	0.491	0.472	0.448
Pair removed, source initialization	0.551	0.529	0.508	0.478	0.551	0.532	0.514	0.488
Independent no-pair, source initialization	0.551	0.529	0.508	0.480	0.551	0.532	0.514	0.490
SCEPI adaptation, target only	0.537	0.515	0.489	0.439	0.537	0.519	0.491	0.459
Pair removed, target only	0.549	0.527	0.499	0.459	0.549	0.530	0.502	0.469
Independent no-pair, target only	0.549	0.527	0.499	0.459	0.549	0.530	0.502	0.469
MoCHI, order 1	0.545	0.497	0.453	0.452	0.543	0.499	0.458	0.427
MoCHI, order 2	0.532	0.498	0.453	0.444	0.533	0.497	0.464	0.430
MAVE-NN, order 1	0.476	0.457	0.429	0.407	0.484	0.464	0.431	0.415
MAVE-NN, order 2	0.481	0.460	0.442	0.425	0.484	0.461	0.446	0.437
RFA, order 2	0.544	0.524	0.446	0.398	0.548	0.528	0.441	0.409
Mean ridge	0.749	0.687	0.633	0.576	0.792	0.716	0.651	0.610
Mean MLP	0.800	0.715	0.657	0.604	0.793	0.738	0.673	0.619

Table 18. Phenotype Spearman correlation on common finite records for functional double mutants. Every model has five fitting repetitions. Assay counts are given for each budget. Conditions are averaged within assay before equal-assay averaging. Higher values are better.

	Random combinations				Unseen site pairs			
Predictor	C128	C512	C2048	C-all	C128	C512	C2048	C-all
Assays	5	6	5	6	5	6	5	6
SCEPI adaptation, source	0.739	0.731	0.828	0.771	0.748	0.737	0.827	0.763
Pair removed, source initialization	0.754	0.734	0.830	0.767	0.758	0.742	0.829	0.767
Independent no-pair, source initialization	0.755	0.734	0.830	0.769	0.755	0.745	0.830	0.767
SCEPI adaptation, target only	0.716	0.719	0.824	0.766	0.731	0.734	0.824	0.769
Pair removed, target only	0.720	0.723	0.823	0.767	0.733	0.733	0.824	0.769
Independent no-pair, target only	0.720	0.723	0.823	0.766	0.732	0.733	0.824	0.769
MoCHI, order 1	0.706	0.711	0.824	0.772	0.707	0.718	0.833	0.760
MoCHI, order 2	0.713	0.727	0.836	0.773	0.721	0.738	0.844	0.766
MAVE-NN, order 1	0.726	0.731	0.830	0.775	0.753	0.741	0.838	0.769
MAVE-NN, order 2	0.726	0.732	0.831	0.769	0.751	0.741	0.840	0.763
RFA, order 2	0.705	0.692	0.791	0.747	0.716	0.692	0.798	0.739
Mean ridge	0.602	0.592	0.695	0.662	0.617	0.597	0.691	0.659
Mean MLP	0.557	0.548	0.662	0.649	0.566	0.555	0.661	0.641

ous phenotypes and official test assignments. Each split is evaluated separately across five repetitions. Hydro three-to-many averages its three target backgrounds equally. The background-transfer tests provide training measurements from the other backgrounds. Per-background MoCHI, MAVE-NN, and RFA fits require target-background training labels and are unavailable in those three tests. Pooled representation controls and the adaptation variants use the permitted cross-background training pool.

Tables 19 and 20 report every split. Source-initialized pair fitting improves the TrpB by-position correlation from 0.167 to 0.233. It lowers the one-to-many correlation from 0.252 to −0.214. On Hydro three-to-many, the corresponding values are 0.356 and 0.201. The source-initialized independent no-pair model attains 0.703 on amylase. These results give separate comparisons of source initialization, target fitting, and inference-time pair removal. They establish task-dependent performance for the predicted-single adaptation variant.

Table 19. FLIP2 phenotype Spearman correlation on TrpB, amylase, and IRED. All values use the official test split and average five repetitions. TrpB columns distinguish one-to-many, two-to-many, and by-position protocols. Amylase uses one-to-many; IRED uses two-to-many. Higher values are better.

	TrpB			Amylase	IRED
Predictor	1-to-many	2-to-many	By position	1-to-many	2-to-many
SCEPI adaptation, source	−0.214	0.155	0.233	0.697	0.200
Pair removed, source initialization	0.252	0.172	0.167	0.697	0.221
Independent no-pair, source initialization	0.252	0.215	0.172	0.703	0.221
SCEPI adaptation, target only	0.388	0.206	0.180	0.677	0.222
Pair removed, target only	0.388	0.413	0.180	0.677	0.219
Independent no-pair, target only	0.388	0.413	0.180	0.677	0.219
MoCHI, order 1	0.247	0.481	0.256	0.593	0.173
MoCHI, order 2	0.336	0.482	0.218	0.621	0.191
MAVE-NN, order 1	0.245	0.393	0.254	0.574	0.175
MAVE-NN, order 2	0.338	0.460	0.135	0.584	0.175
RFA, order 2	0.319	0.180	0.263	0.625	0.181
Mean ridge	0.454	0.466	0.283	0.693	0.199
Mean MLP	0.437	0.485	0.302	0.695	0.173

J Interpretation of the selection endpoint

Complete, point-valued quartets define the retrospective evaluation cohort. Their double-mutant outcomes are excluded from model training and candidate scores. Measurement status and the quartet margin define evaluation strata. The margin is the minimum of $y+1$ and $5-y$ across all four observed endpoints. A threshold of 0.5 kcal/mol excludes quartets close to either assay endpoint. This analysis conditions the evaluation on measured outcomes inside the assay range.

Recovery is $y_{AB} - \max(y_A, y_B)$. A hit requires at least 0.5 kcal/mol recovery, and both singles must individually lose at least 0.5 kcal/mol relative to wild type. Hit rate and mean recovery have different roles: the first measures successful budgeted discoveries, and the second measures continuous candidate quality. A candidate can meet the recovery criterion while remaining less stable than wild type.

An initial analysis also examined a broader continuous stability pool. On its seed11 cohort of 65 proteins, common-main top-10 mean stability changes differed by −0.0018 kcal/mol from additive ranking and by −0.0054 kcal/mol from global nonlinear ranking. This analysis concerns general stability ranking. The main discovery claim concerns the explicitly defined compensation task.

Table 20. FLIP2 Hydro phenotype Spearman correlation. Three-to-many averages three target backgrounds; each remaining column holds out the indicated background. Every value averages five repetitions. Dashes indicate local fits without target-background training labels. Higher values are better.

Predictor	Three-to-many	P01053	P06241	P0A9X9
SCEPI adaptation, source	0.201	0.368	0.193	0.025
Pair removed, source initialization	0.356	0.420	0.197	0.188
Independent no-pair, source initialization	0.356	0.420	0.215	0.188
SCEPI adaptation, target only	0.059	0.408	0.219	0.036
Pair removed, target only	0.284	0.408	0.251	0.193
Independent no-pair, target only	0.284	0.404	0.258	0.209
MoCHI, order 1	0.549	—	—	—
MoCHI, order 2	0.557	—	—	—
MAVE-NN, order 1	0.554	—	—	—
MAVE-NN, order 2	0.565	—	—	—
RFA, order 2	0.519	—	—	—
Mean ridge	0.586	0.239	0.025	0.168
Mean MLP	0.614	0.317	0.303	0.025

K Random selection and compensatory enrichment

For protein p, let N_p be the number of eligible candidates and H_p the number satisfying the measured recovery threshold. Uniform sampling of K candidates without replacement gives a hit count $X_p \sim \text{Hypergeometric}(N_p, H_p, K)$. Its expected hit rate is $q_p = H_p/N_p$, and

$$\text{Var}(X_p/K) = \frac{q_p(1-q_p)(N_p-K)}{K(N_p-1)}. \tag{16}$$

Expected mean recovery equals the candidate-pool mean. Both expectations are independent of K. They are calculated exactly from every eligible candidate; no random-ranking simulation is needed.

Let $\overline{h}_p(K)$ be a predictor's hit rate averaged over its five fitted repetitions. For the same P eligible proteins, enrichment is

$$E(K) = \frac{|P|^{-1}\sum_{p\in P}\overline{h}_p(K)}{|P|^{-1}\sum_{p\in P} q_p}. \tag{17}$$

This ratio uses the same equal-protein aggregation as the primary results. It remains defined when individual proteins have no observed hits. The standard cohort contains 79,564 candidates and 1,718 hits across 120 proteins. Its pooled candidate hit fraction is 2.16%, and its equal-protein mean random hit rate is 2.63%. The latter is the denominator for every reported enrichment.

This analysis was added after manuscript review. It retains the original model predictions, candidate eligibility, recovery threshold, budgets, and measurement margins. Each predictor's five repetitions are averaged within protein before aggregation. Uncertainty uses 10,000 paired homology-group bootstrap draws with seed 20261019. Each draw resamples model and random results together and recalculates both means and their ratio. These intervals quantify sampling uncertainty in the observed cohort. The exact random expectation has no Monte Carlo ranking error.

At $K = 10$, SCEPI's hit-rate gain over random selection has a 95% interval of [4.01, 12.07] percentage points. Its recovery gain has an interval of [0.499, 0.720] kcal/mol. The stricter homology grouping gives 4.24-fold native enrichment. A 0.5 kcal/mol quartet margin retains 115 proteins and 75,017 candidates. Native enrichment is 4.15 and 4.19 under the 30% and 25% groupings. Common-main-effect enrichment is 4.17 and 4.20 in the same margin cohorts (Table 22).

Table 21. SCEPI native selection relative to exact uniform random selection under 30% homology grouping. The same 120 proteins and 44 groups enter every budget. Random hit rate is 2.63%. Hit-rate gains are percentage points; recovery gains are kcal/mol. Brackets give 95% homology-group bootstrap intervals for enrichment.

Budget	SCEPI hit rate (%)	Hit-rate gain	Enrichment (×)	Recovery gain
1	14.50	11.87	5.52 [2.58, 8.24]	0.784
3	12.89	10.26	4.90 [2.59, 7.11]	0.728
10	11.22	8.59	4.27 [2.74, 5.89]	0.643
30	8.56	5.93	3.26 [2.43, 4.40]	0.529

Table 22. Top-10 random-selection sensitivity. Random rates and model rates are equal-protein means. Native and common-main-effect columns report SCEPI enrichment over the same random expectation. The 30% grouping has 44/44/43 groups at the three margins; the 25% grouping has 39/39/38.

Grouping	Margin	Proteins	Random hit (%)	Native (×)	Common main (×)
30%	0.00	120	2.63	4.27	4.29
30%	0.25	120	2.68	4.17	4.19
30%	0.50	115	2.68	4.15	4.17
25%	0.00	120	2.63	4.24	4.25
25%	0.25	120	2.68	4.17	4.17
25%	0.50	115	2.68	4.19	4.20